SPENDNEST REVENUE INTELLIGENCE WORKSPACE

SLIDE 1: The Manager Problem & 3 Evidence-Backed Findings

Solving information fragmentation, unverified upstream labels, and late-arriving transcripts.

The Manager Challenge

Sales managers receive incoming calls, meetings, and transcripts across disparate systems with unverified labels, missing summaries, and late-arriving recordings. Generic CRMs force every conversation into an artificial pipeline.

Finding 1: Post-Sale Defect Escalation (M07)

Observation: Customer A demonstrated 4 critical billing & reimbursement bugs (event description edit failing to persist, Recreate payment upload failure, stuck processing reimbursements, cached INR 18,200 amount).
Source Reference: Meeting M07, Turns 3, 12, 23, 25.
Product Decision: Implemented multi-route model pairing Customer Follow-up with Sales Coaching/Engineering Escalation; late-transcript update mechanism preserving manual overrides.

Finding 2: Telephony Rep vs Spoken Caller Mismatch (C06)

Observation: Call metadata assigned rep_name as 'Meera', but caller verbally stated 'Ravi here. I'll say at the start that this is a cold call.'
Source Reference: Call C06, Short exact quote.
Product Decision: Displayed both stored metadata and verbal transcript proof side-by-side; added Sales Coaching route to audit telephony account sharing and caller ID profiles.

Finding 3: Regulatory Compliance Boundary (M01)

Observation: Harborpath Advisory assumed SpendNest opens remote personal accounts for overseas individuals; Team member A established strict boundary that SpendNest only serves registered US business entities.
Source Reference: Meeting M01, Turns 2, 4.
Product Decision: Separated Partner Coordination from direct customer sales; mandated compliance boundary verification in the inspection view.

SPENDNEST REVENUE INTELLIGENCE WORKSPACE

SLIDE 2: Interface Decisions & Deliberate Omissions

Prioritizing information clarity, traceability, and verifiable operational signals over vanity charts.

Core Interface Decisions

• Traceable Evidence Display: Every analysis directly cites 1-based turn numbers (meetings) or exact quotes (calls).
• Strict Mathematical Denominators: Dashboard defines exact bases: 44 total records, 40.91% usable text coverage (18/44), and 33.33% source-reported connection rate (10/30).
• Role Attribution Safeguards: Telephony reps are explicitly distinguished from calendar meeting organizers.
• Clickable Navigation: Clicking any metric card instantly filters the 44-record inbox.

What Was Deliberately Left Out & Accepted Trade-Off

• Rejected Alternative: Automatically inferring CRM pipeline currency values and deal stages from spoken conversational numbers.
• Why Rejected: DATA_DICTIONARY strictly prohibits inventing pipeline numbers without authoritative deal tables. Spoken figures represent bugs or exploratory budgets, not confirmed bookings.
• Trade-Off Accepted: Zero speculative revenue forecasting or pipeline charts; 100% verifiable operational intelligence.

SPENDNEST REVENUE INTELLIGENCE WORKSPACE

SLIDE 3: Routing/Folder Design & A Difficult Case (M07)

Transparent multi-route taxonomy, conflict-free export hierarchy, and handling late-arriving transcripts.

Multi-Route Taxonomy & Archive Design

- 6 Standard Categories: Customer Follow-up, Deal / Commercial Next Step, Sales Coaching, Partner / Vendor Coordination, Nurture / No Immediate Action, and Manual Review.
- Multi-Route Flexibility: Conversations can legitimately require multiple destinations (e.g. C03 Customer Follow-up & Deal Next Step; C06 Sales Coaching & Nurture).
- Automated vs Manual Separation: Automated suggestions remain distinct from explicit manager manual overrides.
- Organised Export Archive: Generates real downloadable ZIP with /manifest.json, /manifest.csv, /routing_summary.csv, /findings.json, /brief/, /records/ (44 files), and /analysis/ (18 files).

Difficult Case Study: Record M07

- Initial Delivery (Rev 1): Arrived with recording_status 'recorded' but turns was empty ([]). System correctly assigned Manual Review without fabricating analysis.
- Manual Route Override: Manager inspected M07 and manually set route to Customer Follow-up + Sales Coaching.
- Late-Transcript Update (Rev 2): Delivery 02_update.json supplied Revision 2 with a 27-turn transcript demonstrating 4 production software defects.
- Data Contract Compliance: Whole payload replaced by Rev 2; analysis updated to Fresh; manual route preserved throughout! Subsequent re-import of Rev 1 correctly ignored.

SPENDNEST REVENUE INTELLIGENCE WORKSPACE

SLIDE 4: Validation of a Key Number & Analytical Conclusion

Step-by-step verification against raw source records, envelope payloads, and transcript turns.

Validation of Key Number: Usable Text Coverage

- The Number: Usable text coverage is exactly 18 / 44 (40.91%) post-update (17 / 44 = 38.64% initially).
- Calls Audit: 30 calls in 01_initial.json. Exactly 10 calls (C01-C10) have non-null, non-blank transcript. 20 calls (B01-B20) have transcript=null.
- Meetings Audit: 14 meetings. Exactly 7 meetings (M01-M04, M06, M09, M10) initially contain non-blank speaker turns. 7 meetings (M05, M07, M08, M11-M14) have turns=[]. Initial usable = 10 + 7 = 17.
- Post-Update Verification: In 02_update.json, M07 Rev 2 supplies 27 turns. Total usable becomes 10 + 8 = 18. Distinct records remains strictly 44. Verified by automated tests.

Validation of Analytical Conclusion: C06 Telephony Discrepancy

- The Conclusion: Call C06 represents a high-priority Sales Coaching item for telephony credential audit.
- Source Metadata: In 01_initial.json C06, payload fields record: rep_name: 'Meera', rep_email: 'meera@example.com'.
- Transcript Evidence: Verbatim text begins: 'Good evening. Ravi here. I'll say at the start that this is a cold call.'
- Operational Impact: Grounded inspection caught the identity conflict, proving that previous system tags cannot be treated as ground truth.

SPENDNEST REVENUE INTELLIGENCE WORKSPACE

SLIDE 5: Trade-Offs, AI Usage, Limitations & Next Steps

Transparent engineering assessment, time allocation, and production roadmap.

AI Usage & Deterministic Architecture

• AI Disclosure: Analyses for the 18 usable conversations and the Manager Brief were carefully pre-generated and verified against source transcripts. Deduplication, revisions, filters, and routing execute deterministically at runtime.
• Zero Runtime Hallucination: No live LLM API was invoked at runtime to avoid hallucinating facts or violating local boundaries. No external API keys required.
• Data Contract Compliance: Identity is strictly (source_type, source_id); revision upserts replace whole payloads; manual overrides are persisted in browser localStorage across refresh.
• Time Allocation: Approximately 7.5 hours invested across data contract engineering, UI implementation, automated testing, and presentation artifacts.

Known Limitations & Future Improvements

• Single-Day Window: The dataset captures a curated slice on 21 August 2026 (Asia/Kolkata). It does not represent longitudinal rep quotas.
• No Direct Upstream CTI Sync: Telephony updates currently require batch JSON ingestion rather than a live CTI webhook.
• Future Step 1 (Live CRM Integration): Connect bi-directional sync to Salesforce/HubSpot to write back verified action items.
• Future Step 2 (Private LLM Sidecar): Embed a local model (e.g. Ollama/Llama 3) for real-time draft analysis on new incoming audio transcripts without cloud data leakage.